

Part I.

Reviewed notation

1. Multivariable Regression

- $y : [m \times 1]$ - Output variable (dependent variable)
- $X : [m \times (n + 1)]$ - Input variables (independent variables)
- $\theta : [(n + 1) \times 1]$ - parameters
- m : number of observations
- n : number of parameters
- $h_\theta = X\theta$ = regression output, should be as close as possible to y

1.1. Normal resolution

$$\theta = (\theta^T \theta)^{-1} X^T y \quad (1)$$

1.2. Gradient Descent

Gradient descent (some sort of Newton), used as a numerical method to iteratively find the optimum of a function (J in this case)

$$\theta := \theta - \alpha \Delta J(\theta) \quad (2)$$

$$:= \theta - \alpha \frac{1}{m} X^T (X\theta - y) \quad (3)$$

2. Multivariable regression with Regularization

To automatically choose optimal parameters and avoid from overfitting, an additional feature is introduced λ .

The Normal Method thus becomes

$$\theta = \left(\theta^T \theta + \lambda \begin{pmatrix} 0 & 1 & \dots & 1 \\ \vdots & 1 & \dots & \vdots \\ 0 & 1 & \dots & 1 \end{pmatrix} \right)^{-1} X^T y \quad (4)$$

and the gradient descent algorithm

$$\theta := \theta - \alpha \frac{1}{m} X^T (X\theta - y) + \alpha \frac{\lambda}{m} \theta \quad (5)$$

3. Logistic Regression

3.1. Variables for one output

- $y \in [0, 1]^{m \times 1}$
- $h_\theta = \mathbb{P}(y = 1 | X; \theta) = g(X\theta)$

3.2. Sigmoid

Map $z \in \mathbb{R}$ to $g(z) \in [0, 1]$

$$g(z) = \frac{1}{1 + e^{-z}} \quad (6)$$

3.3. Gradient Descent

$$\theta := \theta - \alpha \frac{1}{m} X^T (g(X\theta) - y) \quad (7)$$

3.4. Variables for p outputs - multi-classifier

- $y \in [0, 1]^{m \times p}$
- $\theta : [(n+1) \times p]$
- $X : [m \times (n+1)]$
- $h_\theta = \begin{pmatrix} h_{\theta_1} \\ \vdots \\ h_{\theta_p} \end{pmatrix}$

4. Neural Networks

4.1. Forward Propagation

- m : number of observations
- $y : [m \times s_L]$
- $\theta^{(l)} : [(s_l + 1) \times s_{(l+1)}]$

Initialization ($l = 1, s_1 = n$):

$$a^{(1)} = x \quad (8)$$

$$[m \times s_1] = [m \times n] \quad (9)$$

$$\tilde{a}^{(1)} = \tilde{x} \quad (10)$$

$$[m \times (s_1 + 1)] = [m \times (n + 1)]$$

$l < L$:

$$\tilde{a}^{(l+1)} = g(a^{(l)} \theta^{(l)}) = g(z^{(l+1)}) \quad (11)$$

$$[m \times (s_l + 1)] = [m \times (s_l + 1)][(s_l + 1) \times s_{(l+1)}] \quad (12)$$

$$a^{(l)} = \begin{pmatrix} 1 \\ \vdots \\ 1 \end{pmatrix} : + \tilde{a}^{(l)} \quad (13)$$

→ add a column of ones at the beginning of the matrix

4.2. Backward Propagation (in order to find ΔJ)

Initialization $l = L$:

$$\begin{aligned}\delta^{(L)} &= a^{(L)} - y^{(L)} \\ [m \times s_L] &= [m \times s_L] - [m \times s_L]\end{aligned}\tag{14}$$

$l < L$

$$\delta^{(l)} = \tilde{\delta}^{(l+1)} \theta^{T(l)} \circ g'(z^{(l)})\tag{15}$$

$$\begin{aligned}&= \tilde{\delta}^{(l+1)} \theta^{T(l)} \circ a^{(l)} \circ (1 - a^{(l)}) \\ [m \times (s_l + 1)] &= [m \times s_{l+1}] - [s_{l+1} \times (s_l + 1)]\end{aligned}\tag{16}$$

4.3. Compute/train a Neural Network

A Neural network is defined by the following dimensions:

- m : number of observations
- s_l : number of node for layer l
- n : number of incoming variables $n = s_1$ ($n + 1$ for vector of ones)
- p : number of output variables ($p = s_L$)
- L : Number of nodes

1. $a^{(1)} = x$
2. forward: get $a = (l)$
3. backward: get $\delta^{(l)}$
4. $\Delta^{(l)} = a^{T(l)} \tilde{\delta}^{(l+1)}$
5. $\nabla^{(l)} J(\theta) = D^{(l)} = \frac{1}{m} \Delta^{(l)}$
6. Optimization to find $\hat{\theta}^{(l)}$: $\min_{\theta} J$

5. Support Vector Machine (SVM)

Also referred to as Large Margin Classifiers.

$$J(\theta) = \min_{\theta} C \sum_{i=1}^m \left[y^{(i)} \text{cost}_1(\theta^T x^i) + (1 - y^{(i)}) \text{cost}_0(\theta^T x^{(i)}) \right] + \frac{1}{2} \sum_{j=1}^n \theta_j^2.\tag{17}$$

$$C = \frac{1}{\lambda}\tag{18}$$

Part II.

First Version - as in Coursera Course

Cost Function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2 \quad (19)$$

Hypothesis

$$h_{\theta}(\mathbf{x}) = \theta^T \mathbf{x} \quad (20)$$

with $\mathbf{x} = (1, x_1, \dots, x_n)$ and $\theta = (\theta_0, \theta_1, \dots, \theta_n)$

Gradient Descent Minimize $J(\theta)$ for θ , the algorithm is

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \quad (21)$$

where α is the *learning rate*.

and

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \quad (22)$$

alternatives to gradient descent

- Conjugate gradient
- BFGS
- L-BFGS

Analytical Solution

$$\theta = (X^T X)^{-1} X^T y \quad (23)$$

6. Logistic Regression

classification: $y = 0$ or $y = 1$

Want : $0 \leq h_{\theta}(x) \leq 1$

$$h_{\theta}(x) = g(\theta^T x) \quad (24)$$

with logistic (sigmoid) function $z \in \mathbb{R}$:

$$g(z) = \frac{1}{1 + e^{-z}} \in (0, 1) \quad (25)$$

output can be read as $h_{\theta}(x) = \mathbb{P}(y = 1|x; \theta)$: probability that $y = 1$ given x parametrized by θ is h .

6.1. Decision Boundary

$$y = \begin{cases} 1 & z > 0 \\ 0 & z < 0 \end{cases} \quad (26)$$

where z is the argumen of g and is usally of the form $\theta^T x$

$$\text{Cost}(h_\theta(x, y)) = \frac{1}{2}(h_\theta(x) - y)^2 \quad (27)$$

6.2. Cost function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_\theta(x^{(i)}), y^{(i)}) \quad (28)$$

$$\text{Cost}(h_\theta(x, y)) = \begin{cases} -\log(h_\theta(x)) & y = 1 \\ -\log(1 - h_\theta(x)) & y = 0 \end{cases} \quad (29)$$

since $y = 1$ or $y = 0$:

$$\text{Cost}(h_\theta(x), y) = -y \log(h_\theta(x)) - (1 - y) \log(1 - h_\theta(x)) \quad (30)$$

	y	$h_\theta(\mathbf{x})$	$\text{Cost}(h, y)$
	0	0	0
allows to have	1	0	∞
	0	1	∞
	1	1	0

7. Regularization

7.1. Problem of overfitting

If too many features (θ), the learned hypothesis may fit the training examples very well ($J(\theta) \approx 0$), but fail to generalize to new examples.

7.1.1. Addressing overfitting

- Reduce the number of features
 - Manually select features to be kept
 - Model selection algorithm
- Regularization

7.2. Cost function

add λ to reduce the number of allowed parameters.

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_\theta(x^i) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right] \quad (31)$$

If λ is too large, **underfit** occurs.

7.3. Regularized Logistic regression

$$J(\theta) = -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log(h_{\Theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\Theta}(x^{(i)})) \right] + \frac{\lambda}{2m} \sum_{j=1}^m \Theta_j^2. \quad (32)$$

The partial derivative are

$$\frac{\partial J(\theta)}{\partial \theta_0} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \quad \text{for } j = 0 \quad (33)$$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \left(\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \right) + \frac{\lambda}{m} \theta_j \quad \text{for } j \geq 1. \quad (34)$$

8. Neural Networks: Representation

8.1. Non-linear hypotheses

8.2. Neurons and the brain

8.3. Model representation I

Neuron model: Logistic unit

$\mathbf{x} = (1, x_1, \dots, x_n)$ and $\theta = (\theta_0, \theta_1, \dots, \theta_n)$

8.4. Feedforward

- Input layer:

$$a^{(1)} = x \quad (35)$$

- Hidden layer(s)

$$z^2 = \Theta^{(1)} a^{(1)} \quad (36)$$

$$a^{(2)} = g(z^{(2)}) \quad (37)$$

$$\dots \quad (38)$$

$$z^{(k+1)} = \Theta^{(k)} a^{(k)} \quad (39)$$

$$a^{(k+1)} = g(z^{(k+1)}) \quad (40)$$

- Output layer

$$z^{(3)} = \Theta^{(2)} a^{(2)} \quad (41)$$

$$h_{\Theta} = a^{(3)} = g(z^{(3)}) \quad (42)$$

- m total number of sample: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$
- L total number of layers indexed by l
- s_l number of unit per layer (without counting bias unit).
- $\Theta^{(l)}$ has size $s_{l+1} \times (s_l + 1)$ (The $+1$ is here to account for the bias unit).
- Last layer L
 - Binary Classification: $s_L = 1$
 - Multi-class classification: $s_L = K$.

9. Neural networks: learning

9.1. Cost function

$$J(\theta) = -\frac{1}{m} \left[\sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h_{\Theta}(x^{(i)}))_k + (1 - y^{(i)}) \log(1 - h_{\Theta}(x^{(i)}))_k \right] + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\Theta_{ij}^{(l)})^2. \quad (43)$$

9.2. Backpropagation

In order to $\min_{\Theta} J(\Theta)$, need to compute

- $J(\Theta) \Rightarrow$ by plugging values directly into function
- $\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) \Rightarrow$ with backpropagation algorithm

$$\begin{cases} \delta^{(L)} = a^{(L)} - y = h_{\Theta} - y & \text{for } l = L \\ \delta^{(l)} = (\Theta^{(l)})^T \delta^{(l+1)} \circ \frac{d}{dz} g(z^{(l+1)}) & \text{for } 1 < l < L \end{cases} \quad (44)$$

¹ The intuition being δ is the "error" at every node that we try to minimize and converges to the gradient.

Note on $\frac{d}{dz} g(z)$ for $g(z) = \frac{1}{1+e^{-z}}$

$$\frac{d}{dz} g(z) = \frac{e^{-z}}{1 + e^{-z}} = g(z)(1 - g(z)) \quad (46)$$

Backpropagation Algorithm

- For $i = 1$ to m
 - Set $a(1) = x^{(i)}$
 - Perform forward propagation to compute $a^{(l)}$ for $l = 2, \dots, L$
 - Using $y^{(i)}$ compute $\delta^{(L)} = a^{(L)} - y^{(i)}$
 - Compute $\delta^{(L-1)}, \delta^{(L-2)}, \dots, \delta^{(2)}$
 - $\Delta_{ij}^{(l)} := \Delta_{ij}^{(l)} + a_j^{(l)} \delta_i^{(l+1)}$ or $\Delta^{(l)} := \Delta^{(l)} + \delta^{(l+1)} (a^{(l)})^T$
- $D_{ij}^{(l)} := \frac{1}{m} \Delta_{ij}^{(l)} + \lambda \Theta_{ij}^{(l)}$ if $j \neq 0$
- $D_{ij}^{(l)} := \frac{1}{m} \Delta_{ij}^{(l)}$ if $j = 0$

¹ \circ is the Hadamard product and produces a matrix where all elements ij are a multiplication of the element ij of the two input matrices:

$$\begin{pmatrix} a_{11} & \dots & a_{1m} \\ a_{21} & \dots & \vdots \\ \vdots & \dots & \vdots \\ a_{n1} & \dots & a_{nm} \end{pmatrix} \circ \begin{pmatrix} b_{11} & \dots & b_{1m} \\ b_{21} & \dots & \vdots \\ \vdots & \dots & \vdots \\ b_{n1} & \dots & b_{nm} \end{pmatrix} = \begin{pmatrix} a_{11}b_{11} & \dots & a_{1m}b_{1m} \\ a_{21}b_{21} & \dots & \vdots \\ \vdots & \dots & \vdots \\ a_{n1}b_{n1} & \dots & a_{nm}b_{nm} \end{pmatrix} \quad (45)$$

$$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = D_{ij}^{(l)} \quad (47)$$

10. Support Vector Machine (SVM)

$$J(\theta) = \min_{\theta} C \sum_{i=1}^m \left[y^{(i)} \text{cost}_1(\theta^T x^i) + (1 - y^{(i)}) \text{cost}_0(\theta^T x^{(i)}) \right] + \frac{1}{2} \sum_{i=1}^n \theta_j^2. \quad (48)$$

The SVM is a *large margin classifier*